

CS3.301: Operating Systems and Networks


Lecture 9 & 10 - Transport Layer (How It Works!)

INSTRUCTOR: DEEPAK GANGADHARAN

Acknowledgement

- The material in this lecture has been gathered, adapted, and generated from multiple sources, combined with the instructor's own experience and knowledge.
- Sources:
 - Computer Networks, 6th Edition — Tanenbaum, Feamster, and Wetherall
 - Computer Networking: A Top-Down Approach — Kurose and Ross
 - Computer Networking Essentials — YouTube channel
 - Other online sources, duly cited where used

Network Protocol Stack (Recap)



Onto the Transport Layer


Service-to-Service Delivery

- Multiplexing and Demultiplexing
- Addressing scheme: Ports
- Two strategies/protocols that allow this:
 - Transmission Control Protocol (TCP) — favours reliability
 - User Datagram Protocol (UDP) — favours efficiency


Let's Go Back

Two processes want to communicate

- Client sends a request to the server.
- The server provides a reply based on the request made by the client.
- E.g., file transfer to get files.
- E.g., web browsing — send a URL and get a page back.
- How do we transport data from process to process?



The Role of the Operating System



- The software component in the OS that supports network calls is called the protocol stack.
- It provides service primitives — which are nothing but system calls. Some kind of API?

A Question to Consider

Any idea what functionalities should be made available to let a process use the network?

Hint: think of the process API


The Socket API

- A simple abstraction for using the network.
- The network service API used to write essentially all network applications.
- Part of every OS and every language [Berkeley Unix, 1983].
 - Allows user-space applications to interact with the networking subsystem.
- Two services:
 - Streams: reliable (connection-oriented).
 - Datagram: unreliable (connectionless).
- Allows applications to attach to the network at different ports.

The Socket API — Core Functions

Function	Description
socket()	Creates a new socket of a certain type (TCP or UDP) and returns a file descriptor
bind()	Associates the socket with a specific IP address and port
listen()	For server sockets — allows the socket to listen for incoming connections
accept()	For server sockets — waits for a client to connect, then returns a new file descriptor
connect()	For client sockets — initiates a connection to a server
send() / recv()	Transmit data or receive data
close()	Terminates the connection

Using Sockets



More About Ports

- An application process is identified by the tuple (IP address, Protocol, Port).
 - Ports are 16-bit integers representing "mailboxes" that a process leases.
- Servers are often bound to "well-known ports."
- Clients are assigned ephemeral ports.
 - Chosen by the OS temporarily.

Some Well-Known Ports

Port	Protocol	Use
20, 21	FTP	File Transfer
22	SSH	Remote login
25	SMTP	Email
53	DNS	Domain name resolution
80	HTTP	World Wide Web
443	HTTPS	Secured web
554	RTSP	Media player control

An Opportunity for a Context Switch?

- The calls involved in establishing a socket connection are blocking calls:
 - `connect()`, `accept()`, `recv()`
- Once the call is made, the OS halts the program to wait for some response.
- These are essentially system calls.
- A trap instruction is executed, and there is an opportunity for a context switch.

Let Us Take a Step Back

Types of Links

- Full Duplex

- Bidirectional — both sides can send at the same time.

- Half-Duplex


- Bidirectional — but only one direction at a time (e.g., walkie-talkies).

- Simplex


- Unidirectional only (e.g., a keyboard sending input, or a TV broadcast).

Transport Services and Protocols

- Provides logical communication between application processes running on different hosts.
- Transport protocol actions in the end systems:
 - Sender: breaks application messages into segments, passes to the network layer.
 - Receiver: reassembles segments into messages, passes to the application layer.
- Protocols: TCP, UDP




Quick Recap



- Segments carry data across the network.
- Segments are carried within packets, which are carried within frames.
- Each layer adds a header. (Anything above L4 gets bundled and treated as the L4 payload.)


Multiplexing and Demultiplexing

- Multiplexing (as sender): handle data from multiple sockets, add a transport header.
- Demultiplexing (as receiver): use header info to deliver received segments to the correct socket.



Multiplexing and Demultiplexing


- Multiplexing (as sender): handle data from multiple sockets, add a transport header.
- Demultiplexing (as receiver): use header info to deliver received segments to the correct socket.



Each segment stays self-contained — never merged with another process's data

Working of Demultiplexing

- Host receives IP datagrams.
 - Each datagram has a source IP address and a destination IP address.
 - Each datagram carries one transport-layer segment.
 - Each segment has a source port number and a destination port number.
- IP addresses and ports together are used to direct a segment to the appropriate socket.



TCP/UDP Segment Format

Connection-Oriented vs. Connectionless

Demultiplexing Scenarios

■ Connection-Oriented (TCP)

- TCP socket identified by a 4-tuple.
 - Source IP, destination IP, source port, destination port.
- Receiver uses all 4 values to direct a segment to the right socket.
- A server may support many TCP sockets — each socket has its own client.

■ Connectionless (UDP)

- UDP socket identified by a 2-tuple.
 - Destination IP and destination port.
- Receiver uses just the port to redirect to the corresponding socket.
- UDP segments with the same destination port but different source IP/port are redirected to the same socket.

TCP vs UDP

TCP	UDP
Connection Oriented	Not Connection Oriented
Reliability (order maintained, retransmission)	No reliability at L4
Higher overhead — reliability, error checking, etc.	Low overhead
Flow control (based on network conditions)	No implicit flow control
Error detection — retransmits erroneous packets	Some error checking — erroneous packets discarded without notification
Congestion control	No congestion control
Use cases: HTTP/HTTPS, file transfer, mail	Use cases: streaming data, VoIP, DNS queries, ...

Connection-Oriented and Reliability

■ Connection-Oriented

- In TCP, a connection is established first, before any data is transmitted.
- In UDP, there is no notion of a connection starting or ending (relies on timeout instead).

■ Reliability

- TCP confirms data delivery via acknowledgments.
 - Order is preserved / maintained.
 - Errors can be detected and handled (awareness) — TCP can recover from them.
- In UDP there is no confirmation — the sender simply trusts that someone is there to receive the data ("fire and forget").
 - No error awareness at L4 — the protocol itself does not handle it.

Flow Control and Overhead

■ Flow Control

- TCP can adjust its transmission rate to use the maximum available bandwidth.
- It checks how much the receiver can actually receive, and adjusts accordingly.

■ Overhead


- TCP adds a larger header to the data — roughly 20 bytes or even more.
- TCP has more built-in features that simply don't exist in UDP.
- In UDP, the header length is only about 8 bytes.



L4 header + payload

UDP Segment Header

- Length: the size in bytes of the UDP segment, including the header.
- Checksum: for error detection.
 - A 16-bit value computed from UDP header, payload, and a pseudo-header from the IP layer.
 - Supports error detection.
 - Makes use of one's-complement arithmetic to compute the sum.



UDP Segment Format

Checksum Process

■ Sender

- All contents of the header, including IP addresses, are treated as a sequence of 16-bit integers.
- Checksum = addition (one's complement) of the segment's content.

Example

1 1 1 0 0 1 1 0 0 1 1 0 0 1 1 0

1 1 0 1 0 1 0 1 0 1 0 1 0 1 0 1

Add it back 1 1 0 1 1 1 0 1 1 1 0 1 1 1 0 1 1

Sum 1 0 1 1 1 0 1 1 1 0 1 1 1 1 0 0

Checksum 0 1 0 0 0 1 0 0 0 1 0 0 0 0 1 1


■ Receiver

- Computes the checksum of the received content.
- Checks if the received checksum and the header's checksum are equal — if so, no error.
- Otherwise, an error is detected.

Where TCP Fits — Quick Recap

- TCP (Transmission Control Protocol) is one of the two core Transport-layer protocols, alongside UDP.
- It provides a connection-oriented, reliable, in-order, byte-stream service between two processes.
- Recall its four big promises:
 - Reliable delivery — every byte sent eventually arrives, or the sender finds out.
 - In-order delivery — bytes are delivered in exactly the order they were sent.
 - Flow control — the sender never overwhelms the receiver's buffer.
 - Congestion control — the sender never overwhelms the network itself.
- How does TCP actually deliver on all four?


The TCP Segment Structure



- Every TCP segment carries this same header, at least 20 bytes long.
- Notice how much richer this is than UDP's simple 4-field header — every extra field here implements one of TCP's reliability, ordering, or control guarantees.

Header Fields I — Ports, Sequence & Ack Numbers

- Source / Destination Port (16 bits each)
 - Used for multiplexing/demultiplexing to the correct socket, exactly as in UDP.
- Sequence Number (32 bits)
 - The byte-stream position of the FIRST byte of data in this segment — not a segment count, a BYTE count.
- Acknowledgement Number (32 bits)
 - The sequence number of the next byte this side EXPECTS to receive — i.e., "I've correctly received everything before this."
 - Only meaningful when the ACK flag is set.



Header Fields II — The Control Flags

SYN	Synchronize — used to initiate a connection (the handshake).
ACK	Indicates the acknowledgement number field is valid.
FIN	Finish — sender has no more data; used to close a connection gracefully.
RST	Reset — abruptly aborts the connection (e.g., on error).
PSH	Push — asks the receiver to pass data up to the application immediately.
URG	Urgent — marks data as urgent (rarely used in practice today).

Header Fields III — Window, Checksum, Options

- Receive Window (16 bits)

- Used for flow control — tells the sender how many more bytes the receiver is currently willing to accept.

- Checksum (16 bits)


- Error detection over the header, payload, and an IP-layer pseudo-header — same one's-complement technique as UDP.

- Urgent Data Pointer (16 bits)


- Only meaningful when URG is set; points to the end of urgent data.

- Options (variable length)

- Optional extensions — e.g., Maximum Segment Size (MSS), window scaling, timestamps, selective acknowledgment (SACK).



Sequence Numbers in Action — A Worked Example



Round-Trip Time (RTT) Estimation

- TCP needs to know how long to wait before assuming a segment is lost.
 - Too short a timeout -> unnecessary retransmissions. Too long -> slow loss recovery.
- TCP estimates RTT continuously using an Exponentially Weighted Moving Average (EWMA):


$$\text{EstimatedRTT} = (1 - a) \times \text{EstimatedRTT} + a \times \text{SampleRTT} \quad (a = 0.125)$$

$$\text{DevRTT} = (1 - b) \times \text{DevRTT} + b \times |\text{SampleRTT} - \text{EstimatedRTT}| \quad (b = 0.25)$$

$$\text{TimeoutInterval} = \text{EstimatedRTT} + 4 \times \text{DevRTT}$$


SampleRTT = one measured round-trip time. DevRTT estimates how much SampleRTT typically varies, so the timeout adapts to network jitter, not just the average delay.

Reliability via Timeout & Retransmission



Fast Retransmit — Don't Wait for the Timer

- Waiting a full timeout to detect loss is slow.
- Insight: if the receiver gets segments AFTER a gap, it re-sends the SAME ACK for the missing byte — a duplicate ACK.
- Rule: if the sender sees 3 duplicate ACKs (i.e., the same ACK four times total) for the same data, it assumes loss immediately and retransmits — without waiting for the timer.



Flow Control — The Sliding Receive Window


- Purpose: prevent the sender from overwhelming the RECEIVER's buffer.
- The receiver continuously tells the sender, via the Receive Window field, how much free buffer space it currently has.
- The sender is only ever allowed to have this many unacknowledged bytes "in flight."
- As the application reads data out of the receive buffer, the window opens back up and the receiver advertises a larger value.
- If the receiver ever advertises Window = 0, the sender must pause entirely until told otherwise.

Principles of Congestion Control

- Flow control protects the RECEIVER. Congestion control protects the NETWORK.
- Too many senders pushing too much data -> routers' queues overflow -> packets are dropped -> everyone's throughput collapses.
- TCP has no direct way to "ask" the network how much capacity is free.
- Instead, TCP treats packet LOSS as an implicit signal of congestion, and self-limits accordingly.
- Each sender maintains its own congestion window (cwnd): the amount of unacknowledged data it is currently allowed to have in flight, on top of the flow-control window.
 - Actual sending rate is governed by whichever is smaller: cwnd or the receiver's advertised window.

Congestion Control — Slow Start


- At the start of a connection, cwnd begins small (typically 1 MSS).
- For every ACK received, cwnd is increased by 1 MSS.
- Effect: cwnd **DOUBLES** every round-trip time — **exponential growth**.
- Continues until either:
 - A loss event occurs, or
 - cwnd reaches ssthresh (slow-start threshold) — then TCP switches to the more cautious Congestion Avoidance phase.




Congestion Control — Congestion Avoidance (AIMD)

- Once $\text{cwnd} \geq \text{ssthresh}$, TCP switches from exponential to LINEAR growth.
- Additive Increase: increase cwnd by roughly 1 MSS per round-trip time (not per ACK).
- On a loss event (indicated by triple duplicate ACK, i.e., fast retransmit):
 - Multiplicative Decrease: cwnd is cut in HALF, and ssthresh is set to this new, halved value.
- On a timeout (a more severe signal — no ACKs at all):
 - cwnd is reset all the way down to 1 MSS, and TCP re-enters Slow Start.
- This whole strategy is called AIMD — Additive Increase, Multiplicative Decrease.


TCP Connection Establishment: The 3-Way Handshake



TCP Data Transfer: Reliable, In-Order Delivery



TCP Connection Termination: The 4-Way Close



Summary — How TCP Delivers on Its Promises

- Header structure: ports for demultiplexing, sequence/ack numbers for ordering and reliability, flags for connection control, window for flow control, checksum for error detection, options for extensions (MSS, SACK, timestamps...).
- Reliability: byte-level sequence numbers + cumulative ACKs, an adaptively-tuned retransmission timer (EstimatedRTT, DevRTT), and fast retransmit for quick recovery from isolated loss.
- Flow control: the receive window keeps the sender from overrunning the receiver's buffer.
- Congestion control: slow start (exponential ramp-up), congestion avoidance (linear, AIMD growth), and fast recovery keep the sender from overrunning the NETWORK — producing the characteristic sawtooth pattern.

Question: Memory is required for process execution! How does OS handle all the memory requirements?

Thank You!